

User guide

External researcher guide for disclosure-checked remote analysis.

This guide is for an external researcher using the analysis interface from a remote site.

Access

You receive access after project approval. The approval links:

- your identity;
- the project purpose;
- the datasets and tools you may use;
- the output rules for the project.

You open a secure HTTPS page and sign in through the approved identity provider. You do not get shell access, database access, notebooks with raw rows, or data downloads.

What the assistant does

You ask for an aggregate analysis in plain language. The outside interface may use a frontier model to help translate your request into a tool call. The model does not run the analysis and does not see raw data.

The safepod receives the proposed tool call, validates it, runs only approved fixed tools, checks the output for disclosure risk, and returns either a result or a refusal.

Asking questions

Good requests are specific about the outcome, grouping, and filters:

- `mean spend by age band`
- `mean wellbeing by canton`
- `total spend by device os for purchases`
- `count by sex and age band`

The current executable tool is `aggregate_query`: count, mean, and sum over allowlisted datasets.

Future tools such as GLM, regression, and ANOVA will appear in the manifest only after they have fixed schemas and disclosure checks.

Results

Every request returns one of three statuses:

STATUS	MEANING
released	the aggregate passed validation and disclosure checks
redacted	some cells were suppressed, but the remaining aggregate is safe
denied	no data is returned

Released tables include an `n` column. Counts may be rounded. Small cells, dominated cells, and unsafe combinations are suppressed or denied.

Why requests are denied

Common reasons:

- the request names a field outside the allowed catalogue;
- the request asks for row-level data or per-person output;
- the requested breakdown creates groups that are too small;
- a sequence of requests could isolate a person through differencing;
- the requested tool is only planned and not yet executable;
- the project or user is not approved for the request.

The refusal reason is part of the result. Denied requests do not include a table.

What not to ask for

Do not ask for:

- raw rows;
- donor or participant identifiers;
- free-text comments;
- exact timestamps;
- rare subgroups designed to isolate a person;
- repeated variations intended to reconstruct a suppressed value;
- custom code or SQL.

Those requests are blocked by design.

Tool manifest

The interface uses a public manifest that lists available tools, datasets, dimensions, measures, and limits. The manifest helps the outside model plan a valid request, but it does not authorize execution. The safepod validates every request again before anything runs.

The current endpoint is:

```
/api/manifest
```

Session behavior

The system tracks your session budget and previous released aggregates. A later request may be denied even if it looks harmless on its own, because it combines with earlier outputs in a risky way.

Getting better results

- Ask for a broader grouping before asking for a fine-grained grouping.
- Prefer planned project variables and predeclared contrasts.
- Use standard aggregate language: count, mean, total, grouped by, filtered to.
- If a query is redacted, ask for fewer grouping variables or a broader cohort.
- If a query is denied, treat the denial as final unless your project reviewer

changes the approved analysis plan.